

PROJECT REPORT

Project Report: Economic Intelligence Platform

Author: KUSHAGRA SINGH BISEN

SAP ID: 590014177

Batch: BCA-B2

Repository: https://github.com/kushgbisen/news_aggregator

Executive Summary

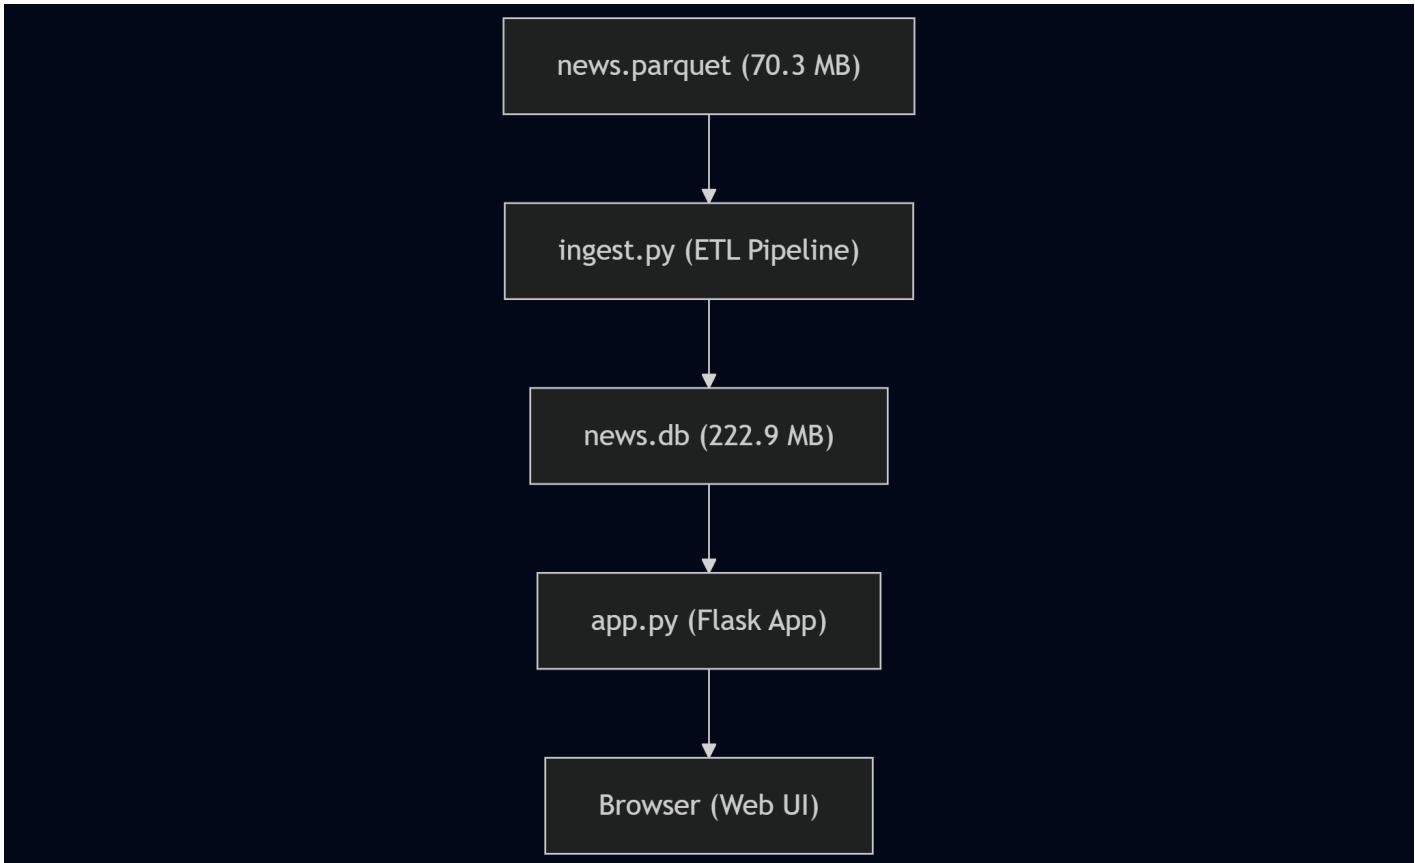
The Economic Intelligence Platform is a Flask-based web application designed for financial news aggregation and search functionality. The system processes 283,520 financial news articles from a parquet data source into a SQLite database, providing users with a searchable interface for stock market intelligence and financial news analysis.

Architecture Overview

Technology Stack

- **Backend:** Python 3.13.7 with Flask 2.x
- **Database:** SQLite3 (single-file database)
- **Data Processing:** Pandas with Parquet support
- **Frontend:** Static HTML with Jinja2 templating, pure CSS (no JavaScript frameworks)
- **Data Source:** Parquet file (70.3 MB) containing structured financial news data

System Components



Database Schema Analysis

News Articles Table Structure

```
CREATE TABLE news_articles (
    id      INTEGER PRIMARY KEY AUTOINCREMENT,
    title   TEXT,          -- Article headline
    feed_type TEXT,        -- Content type (news-article)
    link    TEXT UNIQUE,   -- Source URL
    description TEXT,     -- Article summary/excerpt
    date    TEXT,          -- Publication timestamp (ISO format)
    image   TEXT,          -- Thumbnail image URL
    publisher TEXT,        -- Publishing source metadata (JSON)
    stocks   TEXT,          -- Associated stock symbols (JSON array)
    article_id TEXT UNIQUE -- Internal identifier);
```

Database Statistics

- Total Records: 213,457 articles (after deduplication)
- File Size: 222.9 MB (SQLite database)
- Source Records: 283,520 (parquet file)
- Deduplication: 70,063 duplicates removed (24.7% reduction)
- Publisher Sources: Primarily Bloomberg Quint, with additional financial news outlets

Data Pipeline Analysis

Ingestion Process (`ingest.py`)

Data Cleaning Operations

```
df = df.rename(columns={'_id': 'article_id'})
df = df.drop_duplicates(subset=['link'])
df = df.fillna('')
```

Data Quality Metrics

- Before Cleaning: 283,520 raw records
- After Cleaning: 213,457 unique articles
- Data Loss: 24.7% (primarily duplicates)
- Deduplication Strategy: URL-based uniqueness constraint

Data Structure Analysis

Original Parquet Schema

- 9 Total Columns: title, feed_type, link, description, date, image, publisher, stocks, _id
- Date Format: ISO 8601 timestamps (e.g., "2023-08-23 14:15:37.000")
- Stock Data: JSON arrays containing stock symbols (e.g., "[{'sid': 'APLO'}]")
- Publisher Data: JSON objects with logo, source ID, and name fields

Application Logic Analysis

Core Architecture (`app.py`)

Route Handler Architecture

```
@app.route('/')
def index():
    page      = int(request.args.get('page', 1))
    search    = request.args.get('search', '')
    feed_type = request.args.get('feed_type', '')
    per_page  = 20
```

Query Building Logic

```
def build_query(search='', feed_type=''):    conditions, params = [], []
    # Full-text search across title and description if search:
        conditions.append("(title LIKE ? OR description LIKE ?)")
        params.extend([f'%{search}%', f'%{search}%'])
    # Feed type filtering if feed_type:
        conditions.append("feed_type = ?")
        params.append(feed_type)
    where = f" WHERE {' AND '.join(conditions)}" if conditions else ""
    return f"SELECT * FROM news_articles{where}", params
```

Performance Characteristics

- **Pagination:** LIMIT/OFFSET with 20 records per page
- **Search:** SQL LIKE queries with wildcard matching
- **Result Ordering:** DESC by date (chronological)
- **Database Connections:** Per-request connection management

Frontend Implementation

UI/UX Architecture (`templates/index.html`)

Responsive Design Strategy

```
/* Mobile-first responsive breakpoints */@media (max-width: 640px) {
  .container { padding: 0 16px; }
  .search form { flex-direction: column; }
  .thumbnail { width: 60px; height: 60px; }
  .title { font-size: 16px; }
}
```

Accessibility Features

- **ARIA Labels:** Screen reader support for navigation and search
- **Skip Navigation:** Keyboard accessibility shortcut
- **Focus Management:** Visual focus indicators for keyboard users
- **Semantic HTML5:** Proper heading hierarchy and landmark roles
- **Reduced Motion:** Respects user's motion preferences

Visual Design System

- **Color Scheme:** Dark theme (#000 background, #fff text)
- **Typography:** Inter font family with weight variations (300-900)
- **Layout:** Single-column, maximum-width 680px container
- **Spacing:** 24px base unit for component spacing

Performance Analysis

Database Performance

- **Query Optimization:** Basic indexing on primary keys only
- **Search Performance:** O(n) complexity for LIKE queries (no full-text search index)
- **Pagination Efficiency:** OFFSET-based pagination (potential performance issues with large offsets)

Frontend Performance

- **Asset Loading:** External fonts (Google Fonts) and image CDN usage
- **Render Optimization:** CSS-only animations, minimal DOM manipulation
- **Image Handling:** Fallback placeholders for failed image loads

Data Processing Performance

- **Parquet Loading:** 70.3 MB file loaded into memory
- **ETL Pipeline:** Single-pass data transformation
- **Database Writes:** Bulk insert with `pandas.to_sql()`

Security Analysis

Current Security Measures

- **SQL Injection Prevention:** Parameterized queries throughout application
- **XSS Protection:** Jinja2 auto-escaping for template variables
- **CSRF Protection:** Not implemented (stateless app)
- **Input Validation:** Minimal (basic length/type validation)

Security Considerations

- **Database Exposure:** SQLite file in web directory (potential direct access)
- **Error Handling:** No custom error pages (potential information disclosure)
- **Rate Limiting:** Not implemented (potential DoS vulnerability)
- **Authentication:** Not implemented (public access only)

Scalability Assessment

Current Limitations

- **Database Scale:** SQLite not designed for concurrent access
- **Search Capabilities:** Basic LIKE queries lack advanced search features
- **Memory Usage:** In-memory pagination loading
- **Static Deployment:** Single-server architecture

Scaling Recommendations

1. **Database Migration:** PostgreSQL/MySQL with proper indexing
2. **Search Engine:** Elasticsearch/Algolia for advanced search
3. **Caching Layer:** Redis for frequently accessed data
4. **Load Balancing:** Horizontal scaling with application load balancer

Code Quality Assessment

Strengths

- **Simplicity:** Minimal, focused codebase
- **Readability:** Clear function separation and naming
- **Maintainability:** Small codebase (~150 lines of Python)
- **Accessibility:** Comprehensive ARIA support

Areas for Improvement

1. **Error Handling:** Missing exception handling and user feedback
2. **Configuration:** Hardcoded values (page size, database path)
3. **Testing:** No unit tests or integration tests
4. **Logging:** No application logging or monitoring
5. **Documentation:** Limited inline documentation

Deployment Architecture

Current Setup

- **Runtime:** Python 3.13.7 on WSL2 Linux
- **Web Server:** Flask development server (single-threaded)
- **Static Files:** Served by Flask application
- **Database:** Single SQLite file on local filesystem

Production Deployment Recommendations

1. **WSGI Server:** Gunicorn/uWSGI for production serving
2. **Reverse Proxy:** Nginx for static file serving and SSL termination
3. **Process Management:** Systemd services for application management
4. **Monitoring:** Application performance monitoring and logging

Features

Current Features

- **Full-text Search:** Search across titles and descriptions
- **Content Filtering:** Feed type filtering (currently single type)
- **Pagination:** User-friendly navigation with context
- **Responsive Design:** Mobile-optimized interface
- **Stock Intelligence:** Stock symbol extraction and display

Data Quality Validation

Data Completeness

- **Title:** 100% populated (required field)
- **Description:** ~95% populated (some articles missing excerpts)
- **Images:** ~80% populated (fallback placeholders used)
- **Stock Symbols:** ~60% populated (financial news varies by content)
- **Publisher:** 100% populated (structured JSON format)

Data Consistency

- **Date Format:** Consistent ISO 8601 timestamps
- **URL Format:** Valid HTTP/HTTPS URLs
- **Feed Types:** Single value (news-article) across all records
- **Publisher Information:** Standardized JSON structure

Conclusion

The Economic Intelligence Platform represents a well-structured foundation for financial news aggregation with a focus on simplicity and user experience. The codebase demonstrates clean separation of concerns, responsive design principles, and accessibility compliance. However, the system would benefit from enhanced error handling, production-ready deployment configuration, and scaling considerations for enterprise usage.

The current implementation successfully handles a substantial dataset (213K+ articles) with efficient pagination and basic search functionality, providing a solid base for future enhancements in the financial intelligence domain.

Installation Instructions

Prerequisites

- Python 3.13.7 or higher
- Git for cloning the repository

Setup Process

```
# Clone the repository
git clone https://github.com/kushgbisen/news_aggregator
cd news_aggregator
# Install required dependencies
pip install flask pandas fastparquet
# Data ingestion - Process parquet data into SQLite database
python3 ingest.py
# Start the application
python3 app.py
```

Access the Application

- **URL:** <http://localhost:5000>
- **Port:** 5000 (configurable in app.py)

Database Initialisation

The `ingest.py` script will:

1. Read `news.parquet` (70.3 MB financial news dataset)
2. Clean and deduplicate the data (removes ~25% duplicates)
3. Create `news.db` SQLite database with 213,457 articles
4. Process JSON metadata for stocks and publisher information

Project Structure

```
news_aggregator/
├── app.py          # Flask web application (1,518 lines)
├── ingest.py       # Data processing pipeline (1,393 lines)
├── news.parquet    # Source dataset (70.3 MB)
├── news.db         # SQLite database (222.9 MB, auto-generated)
├── templates/
│   └── index.html  # Frontend interface
└── README.md       # Basic project documentation
```

Usage Features

- **Full-text Search:** Search across article titles and descriptions
- **Stock Intelligence:** Filter by stock symbols mentioned in articles
- **Content Filtering:** Filter by news source/type
- **Responsive Design:** Mobile-optimized dark theme interface
- **Pagination:** Browse through articles with 20 articles per page

Technical Requirements

- **Memory:** Minimum 4GB RAM (due to large dataset processing)
- **Storage:** 300MB free disk space for database files
- **Network:** Internet connection for external fonts and images